

Yuanqing Ma

(+86) 13612923691 | yuanqingma@link.cuhk.edu.cn | github.com/tommyma3 | tommyma3.github.io

EDUCATION

The Chinese University of Hong Kong, Shenzhen

Bachelor of Science in Computer Science and Engineering

Shenzhen, China

Sept. 2024 – Present

- **CGPA:** 3.942/4.000 (Major Ranking: 4/232)
- **Relevant Coursework:** Data Structures, Matrix Computations, Machine Learning, Reinforcement Learning.
- **Research Interests:** Reinforcement Learning, Agentic RL, In-Context RL

RESEARCH EXPERIENCE

The Chinese University of Hong Kong, Shenzhen

Research Assistant

Sept. 2025 – Present

Advised by Prof. Chenjun Xiao

- **Scalable In-Context Reinforcement Learning with Recurrent Algorithm Distillation** | [\[ArXiv\]](#)
Proposed *Recurrent Algorithm Distillation* (RAD), a dual-memory in-context RL agent that recursively compresses long interaction histories into a fixed-size latent memory, while the policy auto-regressively predict actions conditioned on this memory plus a short working memory. Across multiple environments, RAD matches or outperforms AD with up to 4× shorter context and up to 93% fewer inference FLOPs per action.

Tencent

Research Intern

Aug. 2026 – Present

- Developing a label-free framework that enables LLM agents to autonomously generate and refine reusable skills by learning from ordinary deployment traces instead of ground-truth labels or costly online rollouts.

PROJECTS

Self-Summarization Deep Research Agent | [\[GitHub\]](#)

Apr. 2026

- Built a self-summarizing research agent for BrowseComp-Plus that iteratively executes tool calls (e.g. search, retrieve docs) while automatically compacting interaction history to operate within a fixed context window.
- Developed an end-to-end reinforcement learning pipeline with vLLM-based rollout generation, reward-aligned trajectory extraction, and GRPO training using VERL.
- Enabled Qwen3.5-9B to solve long-horizon research tasks with context compaction, achieving a **2× improvement in accuracy** on BrowseComp-Plus.

Deep Reinforcement Learning Implementations (CS185 Course Projects) | [\[GitHub\]](#)

Feb. 2026

- **Imitation Learning:** Implemented Flow Matching vs. MSE for Action Chunking in Push-T environments.
- **Online RL:** Implemented major RL algorithms (e.g. GAE, DQN, SAC) on Atari and Mujoco tasks.
- **LLM RL:** Fine-tuned LLMs using GRPO and GR-REINFORCE for complex mathematical reasoning tasks.
- **Offline RL:** Developed and compared behavior-regularized, value-based, and flow-based offline RL methods (SAC+BC, IQL, FQL) on robotic control environments.

TEACHING

Undergraduate Student Teaching Fellow (USTF)

Sept. 2025 – Present

- **Courses:** Discrete Mathematics (Fall 2026), Linear Algebra (Spring 2026), Intro to CS & Java (Fall 2025).
- Lead weekly tutorial sessions, reinforcing lecture materials through worked examples and practice problems.

HONORS & SCHOLARSHIPS

Academic Performance Scholarship (Class A) & Dean's List: Top-tier academic recognition (AY 2024–25).
Full-tuition Scholarship & Sin Wai Kin Special Admissions Scholarship (2024–Present).

TECHNICAL SKILLS

Languages: Python, Java, C, C++, RISC-V, LaTeX

Libraries: PyTorch, HuggingFace, NumPy, Stable-Baselines3, Transformers, vLLM, SGLang, verl

Tools & Environments: Linux (Ubuntu), Git, GitHub, Gymnasium, Meta-World, OGBench

LEADERSHIP & ACTIVITIES

Shaw College Men's Soccer Team

Sept. 2024 – Present

Assistant Coach & Player

Shenzhen, China

- Gold Medal in the 2025 CUHK(SZ) Olympics and the Shaw Outstanding Student Organization Award.